

PhD Plus Data Literacy in R Cheatsheet

Set working directory

```
setwd("path/to/directory")
Use tab key to drill into directory tree.
Use .. to go back up one branch.
Or Session...Set Working Directory
```

Install/Update/Load Packages

```
install.packages("package")
Or Tools...Install Packages...
library(package)
Tools...Check for Package Updates...
package::function indicates function in package
Example: readr::read_csv()
```

Assignment

Use `<-` or `=`
Alt + - (Win) Option + - (Mac) to insert `<-`

Import data

```
d <- read.csv("path/to/file.csv")
Or using read_csv() from readr package:
d <- read_csv("path/to/file.csv")
Use readxl package to import Excel files.
Use haven package to import SAS, SPSS, Stata files.
```

Glance at data frame named "d"

```
View(d); names(d)
str(d); dplyr::glance(d)
summary(d); head(d); tail(d)
```

Comparison and Logical

```
== (equality)
!= (not equal)
>, >= (greater than, greater than or equal to)
<, <= (less than, less than or equal to)
& (and)
| (or)
! (not)
%in% (matching operator)
```

Missing values

Missing values indicated with NA
NA = not available
`is.na()` returns TRUE if value missing, FALSE otherwise

Create/combine vectors

```
x <- c(2, 4, 8)
y <- c(x, 10) # append 10 to 2,4,8
```

TRUE/FALSE

```
TRUE = 1, FALSE = 0
x <- c(2, 4, 8)
x > 3
[1] FALSE TRUE TRUE
sum(x > 3) # how many TRUE?
[1] 2
```

Basic statistical functions

```
mean(); median(); sd(); var()
quantile() # percentiles
length() # number of values (n)
sqrt() # square root
log() # natural log
log10() # log base 10
min(); max()
range() # min and max
```

Counts and proportions

Count of males/females in column "sex" of data frame "d":
`table(d$sex)`
Proportion of females
`mean(d$sex == "female")`
If any missing values, set `na.rm = TRUE`
`mean(d$sex == "female", na.rm = TRUE)`

2 and 3-way tables

Crosstab of column "sex" (m/f) with column "married" (y/n) in data frame "d":
`xtabs(~ sex + married, data = d)`

Crosstab of column "sex" (m/f) with column "married" (y/n) stratified by "religious" (y/n) in data frame "d":
`xtabs(~ sex + married + religious, data = d)`

marginal proportions

Given following table saved as "tab":

```
      married
sex  n  y
f  20  26
m  30  24
```

Proportion married by sex with base R pipe (`|>`):

```
tab |> proportions(margin = 1)
      married
sex      n      y
f  0.43  0.57
m  0.56  0.44
```

Proportion sex by married with base R pipe (`|>`):

```
tab |> proportions(margin = 2)
      married
sex      n      y
f  0.40  0.52
m  0.60  0.48
```

extract table values

```
tab[1,] # row 1
tab[,2] # column 2
tab[1,2] # cell in row 1, col 2
tab[,1,drop=FALSE] # col 2 as table
```

Plotting with ggplot2

<https://raw.githubusercontent.com/rstudio/cheatsheets/main/data-visualization.pdf>

Example data frame: d

| x | y | g |
|------|-----|-----|
| 1300 | 3.8 | "a" |
| 1400 | 3.2 | "b" |
| 1280 | 2.9 | "a" |

```
library(ggplot2)
```

distribution of y

```
ggplot(d) + aes(x = y) +  
  geom_histogram()  
ggplot(d) + aes(x = y) +  
  geom_density()
```

scatterplot of x and y

```
ggplot(d) + aes(x, y) +  
  geom_point()
```

scatterplot of x and y conditional on g

```
ggplot(d) + aes(x, y) +  
  geom_point() +  
  facet_wrap(~g)
```

scatterplot of x and y points colored by g

```
ggplot(d) + aes(x, y, color=g) +  
  geom_point()
```

scatterplot of x and y, semi-transparent points

```
ggplot(d) + aes(x, y, color=g) +  
  geom_point(alpha = 1/5)  
# alpha ranges from 0  
# (invisible) to 1 (solid)
```

scatterplot with x and y and smooth trend line

```
ggplot(d) + aes(x, y) +  
  geom_point() +  
  geom_smooth()  
# method="lm" for straight line
```

distribution of y for each level of g

```
ggplot(d) + aes(x = g, y = y) +  
  geom_boxplot()  
ggplot(d) + aes(x = g, y = y) +  
  geom_violin()  
ggplot(d) + aes(x = g, y = y) +  
  geom_jitter(width = 0.2,  
              height = 0)
```

Add title, axis labels, etc

```
ggplot(d) +  
  aes(x, y, color=g) +  
  geom_point() +  
  labs(x = "SAT", y = "GPA",  
       title = "SAT vs GPA")
```

Basic data wrangling with dplyr

<https://raw.githubusercontent.com/rstudio/cheatsheets/main/data-transformation.pdf>

Example data frame: d

| x | y | g |
|------|-----|-----|
| 1300 | 3.8 | "a" |
| 1400 | 3.2 | "b" |
| 1280 | 2.9 | "a" |

dplyr functions work with pipes.

Insert pipe operator:

Ctrl+Shift+M (Win); Cmd+Shift+M (Mac)

dplyr always returns a tibble (data frame)

NOTE: Assign result to save transformation!

Extract rows that meet a condition

```
d %>% filter(x > 1300)
```

Arrange data by columns in ascending order

```
d %>% arrange(y)
```

Arrange data by columns in descending order

```
d %>% arrange(desc(y))
```

Select specific columns

```
d %>% select(y, g)  
d %>% select(-x) # all but x
```

Two useful select helpers

```
d %>% select(starts_with("p"))  
d %>% select(-starts_with("p"))  
d %>% select(ends_with("ing"))
```

Add a column and save result

```
d <- d %>%  
  mutate(z = x - mean(x))
```

Summaries for each group (eg, mean)

```
d %>%  
  group_by(g) %>%  
  summarize(m = mean(y))
```

Count membership in group

```
d %>% count(g)
```

Rename columns and save result

```
d <- d %>% rename(SAT = x)  
# new_name = old_name
```

Drop obs missing on a given variable

```
d %>% drop_na(y)
```

Basic data wrangling with dplyr (cont'd)

Create an indicator variable using if_else

```
# j = 1 if y = 4, else 0
d <- d %>% mutate(j =
  if_else(y==4,1,0))
```

Random sample of 20 observations

```
d %>% sample_n(20)
```

Combining dplyr functions and saving result

```
nd <- d %>%
  filter(x > 1000) %>%
  group_by(g) %>%
  summarize(m = mean(y))
```

Working with dates

Use lubridate to format dates. Use m, d, y to create function. Dates stored as number of days since 1/1/70. Eg, to format dates of form May 2, 2021 in column "date" of data frame "d"

```
library(lubridate)
d <- d %>%
  mutate(date = mdy(date))
```

Append hms to format date-times, dates with a time component. Date-times stored as number of seconds since 1/1/70. Eg, to format date-time of form May 2, 2021 2:34:23 in column "date" of data frame "d"

```
d <- d %>%
  mutate(date = mdy_hms(date))
```

Extract day of week, month, year from formatted "date" column in data frame "d":

```
wday(d$date, label = TRUE)
month(d$date, label = TRUE)
year(d$date)
```

Confidence intervals

95% confidence intervals for means

Eg, data frame "d" with column "weight"

```
t.test(d$weight)$conf.int
# or
Hmisc::smean.cl.normal(d$weight)
```

95% confidence intervals for proportions

Eg, data frame "d" with binary column "married" where 1 = married, 0 = not married

```
# proportion married
prop.test(x = sum(d$married),
  n = length(d$married))
```

Linear Models

Model expected value of a variable based on other variables. Eg, data frame "d" with columns "value", "size", "acres", and "zone". Model expected value of value as a function of other variables.

```
m <- lm(value ~ size + acres +
  zone, data = d)
```

View model summary:

```
summary(m)
```

View model coefficients:

```
coef(m)
```

View model diagnostic plots:

```
plot(m)
```

95% confidence intervals for coefficients

```
confint(m)
```

F-test for all coefficients (except intercept)

```
anova(m)
```

lm model summary

- Residuals section: quick assessment of residuals. Ideally 1Q/3Q and Min/Max will be roughly equivalent in absolute value.
- Coefficients: lists the estimated coefficients along with hypothesis tests for the null hypothesis that each coefficient is 0. Est/SE = t-value.
- Residual standard error: estimate of the constant standard deviation of the normal distribution of the errors
- degrees of freedom: sample size - number of coefficients
- R-squared: proportion of variance explained
- F-statistic: overall test that all coefficients (except intercept) are 0.

Visualize interactions in lm model

When two variables interact, their effects depend on each other. Interactions can be visualized with the ggeffects package. Eg, assume model with interaction:

```
m <- lm(value ~ size + acres +
  size:acres, data = d)
```

Visualize interaction with ggffects:

```
# size on x-axis
plot(ggpredict(m, terms =
  c("size", "acres")))
```

```
# acres on x-axis
plot(ggpredict(m, terms =
  c("acres", "size")))
```

Generate citation for R or R package

```
citation() # for R
citation("dplyr") # for dplyr
```